

Introducing `hordes`, a module that makes R available from NodeJS.

About

General philosophy

The general idea of `hordes` is that NodeJS is the perfect tool when it comes to HTTP i/o, hence we can leverage the strength of this ecosystem to build Web Services that can serve R results.

For example, if you have a web service that needs authentication, using `hordes` allows to reuse existing NodeJS modules, which are widely used and tested inside the NodeJS ecosystem. Another good example is NodeJS native cluster mode, and external modules like `pm2` which are designed to launch your app in a multicore mode, watch that your app is still running continuously, and relaunch it if one of the process stop (kind of handy for a production application that handle a lot of load).

`hordes` also makes things easier when it comes to mixing various languages in the same API: for example, you can serve standard HTML on an endpoint, and R on others. And of course, it makes it more straightforward to include R inside an existing NodeJS stack.

From the R point of view, the general idea with `hordes` is that every R function call should be stateless. Keeping this idea in mind, you can build a package where functions are to be considered as ‘endpoints’ which are then called from NodeJS. In other words, there is no “shared-state” between two calls to R—if you want this to happen, you should either register the values inside Node, or use a database as a backend (which should be the preferred solution if you ask me).

But wait...

... Do I need to learn NodeJS?

Yes! How cool is that!

If unlike me you're not a programming language nerd and don't feel like learning JavaScript and NodeJS, the idea is to make the collaboration between R dev and NodeJS developers/production engineer way easier. R devs can write packages with functions returning data that can be handled by NodeJS, and that way, it will be more straightforward to include R inside a web app that already runs on Node, or to build a new Node app that can use the R functions.

Install

`hordes` can be installed from npm with

```
npm install hordes
```

How to

The `hordes` module contains the following functions:

library

`library` behaves as R `library()` function, except that the output is a

JavaScript object with all the functions from the package.

For example, `library("stats")` will return an object with all the functions from `{stats}`. By doing `const stats = library("stats");`, you will have access to all the functions from `{stats}`, for example `stats.lm()`.

Note that if you want to call functions with dot (for example `as.numeric()`), you should do it using the `[]` notation, not the dot one (i.e `base['as.numeric']`, not `base.as.numeric`).

Calling `stats.lm("code")` will launch R, run `stats::lm("code")` and return the output to Node.

Note that every function returns a promise, where R `stderr` rejects the promise and `stdout` resolves it. This point is kind of important to keep in mind if you're building your own package that will be then called through `hordes`.

```
const {library} = require('hordes');
const stats = library(pak = "stats");
stats.lm("Sepal.Length ~ Sepal.Width, data = iris").
  then((e) => console.log(e)).
  catch((err) => console.error(err))
```

Call:

```
stats::lm(formula = Sepal.Length ~ Sepal.Width, data = iris)
```

Coefficients:

```
(Intercept)  Sepal.Width
      6.5262      -0.2234
```

As they are promises, you can use them in an `async/await` pattern or with `then/catch`. The rest of this README will use `async/await`

```
const { library } = require('hordes');
const stats = library("stats");

(async() => {
  try {
    const a = await stats.lm("Sepal.Length ~ Sepal.Width, data = iris")
    console.log(a)
  } catch (e) {
    console.log(e)
  }

  try {
    const a = stats.lm("Sepal.Length ~ Sepal.Width, data = iris")
    const b = stats.lm("Sepal.Length ~ Petal.Width, data = iris")
    const ab = await Promise.all([a, b])
    console.log(ab[0])
    console.log(ab[1])
  } catch (e) {
    console.log(e)
  }
})();
```

Call:

```
stats::lm(formula = Sepal.Length ~ Sepal.Width, data = iris)
```

```
Coefficients:
(Intercept)  Sepal.Width
      6.5262      -0.2234
```

```
Call:
stats::lm(formula = Sepal.Length ~ Sepal.Width, data = iris)
```

```
Coefficients:
(Intercept)  Sepal.Width
      6.5262      -0.2234
```

```
Call:
stats::lm(formula = Sepal.Length ~ Petal.Width, data = iris)
```

```
Coefficients:
(Intercept)  Petal.Width
      4.7776      0.8886
```

Values returned by the `hordes` functions, once in NodeJS, are string values matching the `stdout` of `Rscript`.

If you want to exchange data between R and NodeJS, use an interchangeable format (JSON, arrow, base64 for images, raw strings...):

```
const {library} = require('hordes');
const jsonlite = library("jsonlite");
const base = library("base");

(async () => {
  try {
    const a = await jsonlite.toJSON("iris")
    console.log(JSON.parse(a)[0])
  } catch(e) {
    console.log(e)
  }
  try {
    const b = await base.cat("21")
    console.log(parseInt(b) * 2)
  } catch(e) {
    console.log(e)
  }
}
)();

{
  'Sepal.Length': 5.1,
  'Sepal.Width': 3.5,
  'Petal.Length': 1.4,
  'Petal.Width': 0.2,
  Species: 'setosa'
}
42
```

mlibrary

`mlibrary` does the same job as `library` except the functions are natively memoized.

```
const {library, mlibrary} = require('hordes');
const base = library("base");
const mbase = mlibrary("base");

(async () => {
  try {
    const a = await base.sample("1:100, 5")
    console.log("a:", a)
    const b = await base.sample("1:100, 5")
    console.log("b:", b)
  } catch(e){
    console.log(e)
  }

  try {
    const a = await mbase.sample("1:100, 5")
    console.log("a:", a)
    const b = await mbase.sample("1:100, 5")
    console.log("b:", b)
  } catch(e){
    console.log(e)
  }
})();
```

a: [1] 49 13 37 25 91

b: [1] 5 17 68 26 29

a: [1] 96 17 6 4 75

b: [1] 96 17 6 4 75

get_hash

When calling `library()` or `mlibrary()`, you can specify a hash, which can be compiled with `get_hash`. This hash is computed from the DESCRIPTION of the package called. That way, if ever the DESCRIPTION file changes (version update, or stuff like that...), you can get alerted (app won't launch). Just ignore this param if you don't care about that (but you should in a production setting).

```
const { library, get_hash } = require('hordes');
console.log(get_hash("golem"))

'fdfe0166629045e6ae8f7ada9d9ca821742e8135efec62bc2226cf0811f44ef3'
```

Then if you call `library()` with another hash, the app will fail.

```
var golem = library("golem", hash = "blabla")

      throw new Error("Hash from DESCRIPTION doesn't match specified
hash.")

var golem = library("golem", hash = 'e2167f289a708b2cd3b774dd9d041
b9e4b6d75584b9421185eb8d80ca8af4d8a')
Object.keys(golem).length
```

waiter

You can launch an R process that streams data and wait for a specific output in the stdout.

The promise resolves with and `{proc, raw_output}`: `proc` is the process object created by Node, `raw_output` is the output buffer, that can be turned to string with `.toString()`.

A streaming process here is considered in a loose sense: what we mean here is anything that prints various elements to the console. For example, when you create a new application using the `{golem}` package, the app is ready once this last line is printed to the console. This is exactly what `waiter` does, it waits for this last line to be printed to the R stdout before resolving.

```
> golem::create_golem('pouet')
-- Checking package name -----
v Valid package name
-- Creating dir -----
v Created package directory
-- Copying package skeleton -----
v Copied app skeleton
-- Setting the default config -----
v Configured app
-- Done -----
A new golem named pouet was created at /private/tmp/pouet .
To continue working on your app, start editing the 01_start.R file.
```

```
const { waiter } = require("hordes")
const express = require('express');
const app = express();

app.get('/creategolem', async(req, res) => {
  try {
    await waiter("golem::create_golem('pouet')", solve_on = "To continue
working on your app");
    res.send("Created ")
  } catch (e) {
    console.log(e)
    res.status(500).send("Error creating the golem project")
  }
})

app.listen(2811, function() {
  console.log('Example app listening on port 2811!')
})
```

-> <http://localhost:2811/creategolem>

Changing the process that runs R

By default, the R code is launched by `RScript`, but you can specify another (for example if you need another version of R):


```
Call:
stats::lm(formula = Sepal.Length ~ Sepal.Width, data = iris)
```

```
Coefficients:
(Intercept)  Sepal.Width
      6.5262      -0.2234
```

API using express

```
const express = require('express');
const { library } = require('hordes');
const app = express();
const stats = library("stats");

app.get('/lm', async(req, res) => {
  try {
    const output = await stats.lm(`${req.query.left} ~ ${req.query.right}`)
    res.send('

' + output + '

') } catch (e) { res.status(500).send(e) }) app.get('/rnorm', async(req, res) => { try { const output = await
stats.rnorm(req.query.left) res.send('

' + output + ") } catch (e) { res.status(500).send(e) }) app.listen(2811, function() { console.log('Example app
listening on port 2811!') })
```

->

[http://localhost:2811/lm?left=iris\\$Sepal.Length&right=iris\\$Petal.Length](http://localhost:2811/lm?left=iris$Sepal.Length&right=iris$Petal.Length)

-> <http://localhost:2811/rnorm?left=10>

golem Creator

```
const { waiter } = require("hordes")
const express = require('express');
const app = express();

app.get('/creategolem', async(req, res) => {
  try {
    await waiter(`golem::create_golem('${req.query.name}')`, solve_on = "To
continue working on your app");
    res.send("Created ")
  } catch (e) {
    console.log(e)
    res.status(500).send("Error creating the golem project")
  }
})

app.listen(2811, function() {
  console.log('Example app listening on port 2811!')
})
```

-> <http://localhost:2811/creategolem?name=coucou>